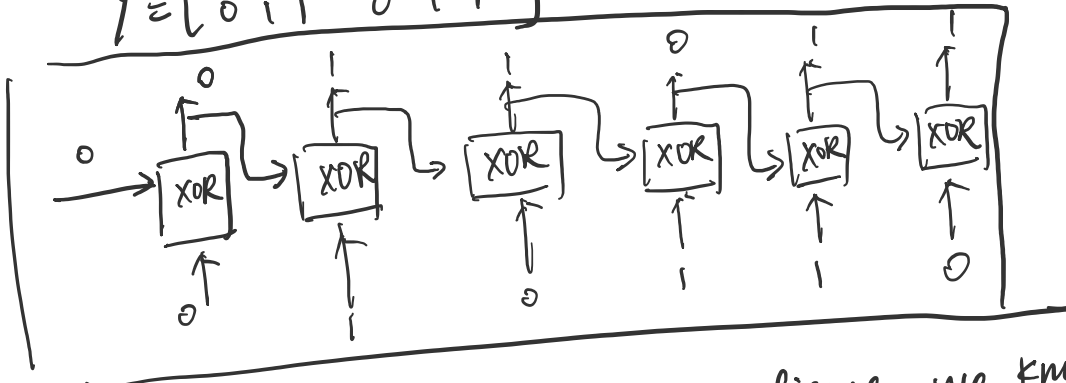


Q1  $x = [0 \ 1 \ 0 \ 1 \ 1 \ 0]$

$y = [0 \ 1 \ 1 \ 0 \ 1 \ 1]$



$y_{t-1}$	$x_t$	$y_t$
0	0	0
0	1	1
1	0	1
1	1	0

From this figure, we know it is  
an XOR operator

$$\begin{cases} y_t = y_{t-1} \bar{x}_t + \bar{y}_{t-1} x_t \\ y_0 = 0 \end{cases}$$

2.  $y_t = h_t$

$$\begin{aligned} h_t &= h_{t-1}(1-x_t) + (1-h_{t-1})x_t \\ &= h_{t-1} - h_{t-1}x_t + x_t - x_t h_{t-1} \\ &= h_{t-1} + x_t - 2h_{t-1}x_t \end{aligned}$$

$h_0 = 0.$

## Q2 LSTM

$$(4): C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

we assume  $C_t = h_t$   $\tilde{C}_t = \bar{h}_{t-1}$ ,  $f_t = \bar{x}_t$ ,  $i_t = x_t$ .

so (4):  $h_t = \bar{x}_t h_{t-1} + x_t \bar{h}_{t-1}$

so (2):  $i_t = \sigma(W_i[h_{t-1}, x_t] + b_i) = x_t$

$$\therefore W_i = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad b_i = 0$$

(1):  $f_t = \sigma(W_f[h_{t-1}, x_t] + b_f) = \bar{x}_t = 1 - x_t$

$$\therefore W_f = \begin{bmatrix} 0 \\ -1 \end{bmatrix} \quad b_f = 1$$

(3):  $\tilde{C}_t = \tanh(W_c[h_{t-1}, x_t] + b_c) = 1 - h_{t-1}$

$$\therefore W_c = \begin{bmatrix} -1 \\ 0 \end{bmatrix} \quad b_c = 1$$

(5):  $O_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$

we make  $O_t \equiv 1$ .

$$\text{so } W_o = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \quad b_o = 1$$

(6):  $h_t = O_t * \tanh(C_t) = \tanh(h_t) \Rightarrow h_t = \tanh(h_t)$   
that make sense.

so the assumption we make is right.

in all.

$$W_f = \begin{bmatrix} 0 \\ -1 \end{bmatrix}$$

$$b_f = 1$$

$$W_i = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$b_i = 0$$

$$W_c = \begin{bmatrix} -1 \\ 0 \end{bmatrix}$$

$$b_c = 1$$

$$W_o = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$b_o = 1$$

Problem 3.

$$\log p(y_t | x, y_{<t}) \leq 0.$$

$$\therefore \max\{y \in B_{t+1}\} \leq \max\{y \in B_t\}$$

hence if  $\max\{y \in B_t\} \leq \text{best}_t$ ,  $\max\{y \in B_{t+1}\} \leq \text{best}_t$

As a result, future steps will be no better and current highest scoring beam is the overall highest probability completed beam.

Q4.

$$h_T = (W^T)^t h_0 \quad (1)$$

$W \in \mathbb{C}^{n \times n}$ , so  $W = V \Lambda V^T$ ,  $V$  is eigen vectors,  
 $\Lambda$  is eigen values

$$h_T = \left[ (V \Lambda V^T)_1 (V \Lambda V^T)_2 \dots (V \Lambda V^T)_t \right] h_0$$

since  $V^T V = E$

$$h_T = [V \Lambda^t V^T] h_0$$

$$\frac{dh_T}{dh_0} = V \Lambda^t V^T$$

if  $t \gg 0$ , if  $\rho(W) < 1$ , then  $\Lambda^t \rightarrow 0$ , resulting  
in vanishing gradient

if  $\rho(W) > 1$   $\Lambda^t \rightarrow \infty$ , resulting in  
exploding gradient.

$$Q5 \quad (a) \quad \text{Agg}(H'_t) = \sum_{j=1}^{\#N(v_i)} f_{ji}(h_j^t)$$

$$h_i^{t+1} = g(h_i^t, \sum_{j=1}^{\#N(v_i)} f_{ji}(h_j^t))$$

$$(b). \quad \text{Agg}(H'_t) = [0.6 \ 0.2 \ 0.2] \begin{bmatrix} f(h_2^t) \\ f(h_3^t) \\ f(h_4^t) \end{bmatrix} = [0.6 \ 0.2 \ 0.2] \begin{bmatrix} [-2 \ 2] \\ [0 \ -2] \\ [2 \ 0] \end{bmatrix}$$


$$= 0.6 [-2 \ 2] + 0.2 [0 \ -2] + 0.2 [2 \ 0]$$

$$= [-0.8 \ 0.8]$$

$$h_i^{t+1} = g([1 \ -1], [-0.4 \ 0.4]) = w(h_i^t)^T + \max([-0.8 \ 0.8], 0)$$

$$= [1 \ 1] \begin{bmatrix} 1 \\ -1 \end{bmatrix} + [0 \ 0.8] = [0 \ 0.8]$$

(U) I think every token has connection to each other . so it would be like



so it has 12 edges


$$(d) \quad \text{Agg}(H'_{it}) = \text{softmax}_j (Q^l h_i^l \cdot K^l h_j^l)$$

$$h_i^{l+1} = g(h_i^l, \text{Agg}(H'_{it})) = \sum_{j \in S} (V^l h_j^l \cdot \text{Agg}(H'_{it}))$$

From above, we see equation 15 has the same format as (17). So the Transformer model's single-head attention mechanism is equivalent to a special case of a GNN.

From (c) For GNN to represent a sentence, it is fully connected. AND the fully connected graph doesn't show any information about sequence order.

For example.



From this figure, we don't know if the original sequence is "I eat Red apple", or "Red apple I eat".



### **Key contributions**

1. Propose a new model(PJ-X) that can: 1. Generate visual and textual explanations 2. Localize salient regions while generating textual rationales.
2. Quantitatively show that training with the textual explanation not only yields better textual justification models, but also better localizes the evidence that supports the decision.
3. Quantitatively show cases where visual explanation is more insightful than textual explanation and vice versa. And multimodal explanation models offer significant benefits over unimodal approaches.
4. Present two novel datasets of human annotated multimodal explanations for activity recognition and visual question answering.

### **Strengths**

1. Well organized and well written.
2. All figures and tables in the essay helps me better understand the paper. Also experiments results are clear and persuasive.

### **Weaknesses**

1. The author writes: sometimes the concept is easily conveyed when looking at the visual pointing result, but sometimes textual justification captures the rationale. While the author doesn't point out or guessed the factor for the deviation of visual or textual explanation.
2. The experiment result is predictable. More information can get better performance. Training with the textual explanation can yield better result.

### Personal takeaway

1. In model training part of this essay, the author uses small dataset to do pretraining on answering model, then freeze the answering model and train explanation model on a smaller dataset. These techniques help to cut down the computational expense.
2. Figure pairs inserted in this essay are often similar in some way, while the textual and visual explanations are different. This makes readers believe the work is persuasive.

# Assignment 4

Your name:

Your GTID:

# Seq2Seq Results [4 points]

- Put your loss & perplexities from training here, both before and after hyperparameter tuning.

Before hyperparameter tuning: Training Loss: 5.1886. Validation Loss: 5.1310.  
Training Perplexity: 179.2245. Validation Perplexity: 169.1792.

After hyperparameter tuning: Training Loss: 4.4945. Validation Loss: 4.4542.  
Training Perplexity: 89.5250. Validation Perplexity: 85.9847.

- Explain what you did here as well, you can use another slide if needed
- I use LSTM, which performs better than RNN. Then I add a weight decay ( $1e-5$ ) parameter to optimizer to prevent overfit. Also I set epoch to 30.

# Transformer Results [5 points]

- Put your loss & perplexities from training here, both before and after hyperparameter tuning.
- Before: Training Loss: 4.9300. Validation Loss: 4.8563. Training Perplexity: 138.3757. Validation Perplexity: 128.5462.
- After: Training Loss: 2.9833. Validation Loss: 3.3124. Training Perplexity: 19.7528. Validation Perplexity: 27.4496.
- Explain what you did here as well, you can use another slide if needed
- I change the learning rate to  $1e-2$ .

# Machine Translation with Seq2Seq and Transformers

In this exercise you will implement a [Sequence to Sequence\(Seq2Seq\)](#) and a [Transformer](#) model and use them to perform machine translation.

**A quick note: if you receive the following `TypeError "super(type, obj): obj must be an instance or subtype of type"`, try re-importing that part or restarting your kernel and re-running all cells.** Once you have finished making changes to the model constructor, you can avoid this issue by commenting out all of the model instantiations after the first (e.g. lines starting with "model = TransformerTranslator(\*args, \*\*kwargs)").

```
In [1]: import numpy as np
import csv
import torch

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

## 1: Introduction

### Multi30K: Multilingual English-German Image Descriptions

[Multi30K](#) is a dataset for machine translation tasks. It is a multilingual corpus containing English sentences and their German translation. In total it contains 31014 sentences(29000 for training, 1014 for validation, and 1000 for testing). As one example:

En: Two young, White males are outside near many bushes.

De: Zwei junge weiße Männer sind im Freien in der Nähe vieler Büsche.

You can read more info about the dataset [here](#). The following parts of this assignment will be based on this dataset.

### TorchText: A PyTorch Toolkit for Text Dataset and NLP Tasks

[TorchText](#) is a PyTorch package that consists of data processing utilities and popular datasets for natural language. The key idea of TorchText is that datasets can be organized in *Field*, *TranslationDataset*, and *BucketIterator* classes. They serve to help with data splitting and loading, token encoding, sequence padding, etc. You don't need to know about how TorchText works in detail, but you might want to know about why those classes are needed and what operations are necessary for machine translation. These knowledge can be migrated to all sequential data modeling. In the following parts, we will provide you with some code to help you understand.

You can refer to torchtext's documentation(v0.6.0) [here](#).

## Spacy

Spacy is package designed for tokenization in many languages. Tokenization is a process of splitting raw text data into lists of tokens that can be further processed. Since TorchText only provides tokenizer for English, we will be using Spacy for our assignment.

**Notice: For the following assignment, we strongly recommend you to work in a virtual python environment. We recommend Anaconda, a powerful environment control tool. You can download it [here](#).**

### 1.1: Prerequisites

Before you start this assignment, you need to have all required packages installed either on the terminal you are using, or in the virtual environment. Please make sure you have the following package installed:

PyTorch, TorchText, Spacy, Tqdm, Numpy

You can first check using either `pip freeze` in terminal or `conda list` in conda environment. Then run the following code block to make sure they can be imported.

```
In [5]: # Just run this block. Please do not modify the following code.
import math
import time

# Pytorch package
import torch
import torch.nn as nn
import torch.optim as optim

# Torchtest package
from torchtext.datasets import Multi30k
from torchtext.data import Field, BucketIterator

# Tqdm progress bar
from tqdm import tqdm_notebook, tqdm

# Code provide to you for training and evaluation
from utils import train, evaluate, set_seed_nb, unit_test_values
```

Once you properly import the above packages, you can proceed to download Spacy English and German tokenizers by running the following command in your **terminal**. They will take some time.

```
python -m spacy download en
```

```
python -m spacy download de
```

Now lets check your GPU availability and load some sanity checkers. By default you should be using your gpu for this assignment if you have one available.

```
In [6]: # Check device availability
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print("You are using device: %s" % device)
```

You are using device: cuda

```
In [7]: # Load checkers
d1 = torch.load('./data/d1.pt')
d2 = torch.load('./data/d2.pt')
d3 = torch.load('./data/d3.pt')
d4 = torch.load('./data/d4.pt')
```

## 1.2: Preprocess Data

With TorchText and Spacy tokenizers ready, you can now prepare the data using *TorchText* objects. Just run the following code blocks. Read the comment and try to understand what they are for.

```
In [9]: # You don't need to modify any code in this block

# Define the maximum length of the sentence. Shorter sentences will be padded to the
MAX_LEN = 20

# Define the source and target language
SRC = Field(tokenize = "spacy",
            tokenizer_language="de_core_news_sm",
            init_token = '<sos>',
            eos_token = '<eos>',
            fix_length = MAX_LEN, lower = True)

TRG = Field(tokenize = "spacy",
            tokenizer_language="en_core_web_sm",
            init_token = '<sos>',
            eos_token = '<eos>',
            fix_length = MAX_LEN,
            lower = True)

# Download and split the data. It should take some time
train_data, valid_data, test_data = Multi30k.splits(exts = ('.de', '.en'),
                                                    fields = (SRC, TRG))
```

/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Field class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.

warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.'.format(self.\_\_class\_\_.\_\_name\_\_), UserWarning)

/usr/local/lib/python3.7/dist-packages/torchtext/data/field.py:150: UserWarning: Field class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.

warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.'.format(self.\_\_class\_\_.\_\_name\_\_), UserWarning)

/usr/local/lib/python3.7/dist-packages/torchtext/data/example.py:78: UserWarning: Example class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.

warnings.warn('Example class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.', UserWarning)

```
In [10]: # Define Batchsize
BATCH_SIZE = 128

# Build the vocabulary associated with each language
SRC.build_vocab(train_data, min_freq = 2)
TRG.build_vocab(train_data, min_freq = 2)
```



```
# Get the padding index to be ignored later in loss calculation
PAD_IDX = TRG.vocab.stoi['<pad>']

# Get data-loaders using BucketIterator
train_loader, valid_loader, test_loader = BucketIterator.splits(
    (train_data, valid_data, test_data),
    batch_size = BATCH_SIZE, device = device)

# Get the input and the output sizes for model
input_size = len(SRC.vocab)
output_size = len(TRG.vocab)
```

/usr/local/lib/python3.7/dist-packages/torchtext/data/iterator.py:48: UserWarning: BucketIterator class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.

warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.'.format(self.\_\_class\_\_.\_\_name\_\_), UserWarning)

## 2: Implement Vanilla RNN and LSTM

In this section, you will need to implement a Vanilla RNN and an LSTM unit using PyTorch Linear layers and nn.Parameter. This is designed to help you to understand how they work behind the scene. The code you will be working with is in *LSTM.py* and *RNN.py* under *naive* folder. Please refer to instructions among this notebook and those files.

### 2.1: Implement an RNN Unit

In this section you will be using PyTorch Linear layers and activations to implement a vanilla RNN unit. You are only implementing an RNN cell unit over one time step! The test case reflects this by having only one sequence. Please refer to the following structure and complete the code in *RNN.py*:



Run the following block to check your implementation

```
In [11]: from models.naive.RNN import VanillaRNN

set_seed_nb()
x1,x2 = (1,4), (-1,2)
h1,h2 = (-1,2,0,4), (0,1,3,-1)
batch = 4
x = torch.FloatTensor(np.linspace(x1,x2,batch))
h = torch.FloatTensor(np.linspace(h1,h2,batch))
rnn = VanillaRNN(x.shape[-1], h.shape[-1], 3)
out, hidden = rnn.forward(x,h)

expected_out, expected_hidden = unit_test_values('rnn')

if out is not None:
    print('Close to out: ', expected_out.allclose(out, atol=1e-4))
    print('Close to hidden: ', expected_hidden.allclose(hidden, atol=1e-4))
else:
    print("NOT IMPLEMENTED")
```

Close to out: True  
Close to hidden: True

/content/drive/My Drive/Colab Notebooks/hw4-code/models/naive/RNN.py:66: UserWarning: Implicit dimension choice for log\_softmax has been deprecated. Change the call to

```
include dim=X as an argument.
output[t] = self.softmax(self.wy(cat))
```

## 2.2: Implement an LSTM Unit

In this section you will be using PyTorch `nn.Parameter` and activations to implement an LSTM unit. You can simply translate the following equations using `nn.Parameter` and PyTorch activation functions to build an LSTM from scratch:

$$\begin{aligned} i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi}) \\ f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf}) \\ g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg}) \\ o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho}) \\ c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\ h_t &= o_t \odot \tanh(c_t) \end{aligned}$$

Here's a great visualization of the above equation from [Colah's blog](#) to help you understand LSTM unit. You can also read more about it from that blog.



If you want to see `nn.Parameter` in example, check out this [tutorial](#) from PyTorch. Run the following block to check your implementation.

Note that in this case we are implementing a full loop with LSTM, iterating over each time step. The test cases reflect this as there are multiple sequences.

```
In [12]: from models.naive.LSTM import LSTM

set_seed_nb()
x1,x2 = np.mgrid[-1:3:3j, -1:4:2j]
h1,h2 = np.mgrid[-2:2:3j, 1:3:4j]
batch = 4
x = torch.FloatTensor(np.linspace(x1,x2,batch))
h = torch.FloatTensor(np.linspace(h1,h2,batch))
lstm = LSTM(x.shape[-1], h.shape[-1])
h_t, c_t = lstm.forward(x)

expected_ht, expected_ct = unit_test_values('lstm')
```

```

if h_t is not None:
    print('Close to h_t: ', expected_ht.allclose(h_t, atol=1e-4))
    print('Close to c_t; ', expected_ct.allclose(c_t, atol=1e-4))
else:
    print("NOT IMPLEMENTED")

```

Close to h\_t: True  
Close to c\_t; True

## 3: Train a Seq2Seq Model

In this section, you will be working on implementing a simple Seq2Seq model. You will first implement an Encoder and a Decoder, and then join them together with a Seq2Seq architecture. You will need to complete the code in *Decoder.py*, *Encoder.py*, and *Seq2Seq.py* under *seq2seq* folder. Please refer to the instructions in those files.

### 3.1: Implement the Encoder

In this section you will be implementing an RNN/LSTM based encoder to model English texts. Please refer to the instructions in *seq2seq/Encoder.py*. Run the following block to check your implementation.

```

In [13]: from models.seq2seq.Encoder import Encoder

set_seed_nb()
i, n, h = 10, 4, 2

encoder = Encoder(i, n, h, h)
x_array = np.random.rand(5,1) * 10
x = torch.LongTensor(x_array)
out, hidden = encoder.forward(x)

expected_out, expected_hidden = unit_test_values('encoder')

print('Close to out: ', expected_out.allclose(out, atol=1e-4))
print('Close to hidden: ', expected_hidden.allclose(hidden, atol=1e-4))

```

Close to out: True  
Close to hidden: True

### 3.2: Implement the Decoder

In this section you will be implementing an RNN/LSTM based decoder to model German texts. Please refer to the instructions in *seq2seq/Decoder.py*. Run the following block to check your implementation.

```

In [14]: from models.seq2seq.Decoder import Decoder

set_seed_nb()
i, n, h = 10, 2, 2
decoder = Decoder(h, n, n, i)
x_array = np.random.rand(5, 1) * 10
x = torch.LongTensor(x_array)
_, enc_hidden = unit_test_values('encoder')
out, hidden = decoder.forward(x, enc_hidden)

expected_out, expected_hidden = unit_test_values('decoder')

```

```
print('Close to out: ', expected_out.allclose(out, atol=1e-4))
print('Close to hidden: ', expected_hidden.allclose(hidden, atol=1e-4))
```

```
Close to out: True
Close to hidden: True
```

### 3.3: Implement the Seq2Seq

In this section you will be implementing the Seq2Seq model that utilizes the Encoder and Decoder you implemented. Please refer to the instructions in `seq2seq/Seq2Seq.py`. Run the following block to check your implementation.

```
In [15]: from models.seq2seq.Seq2Seq import Seq2Seq

set_seed_nb()
embedding_size = 32
hidden_size = 32
input_size = 8
output_size = 8
batch, seq = 1, 2

encoder = Encoder(input_size, embedding_size, hidden_size, hidden_size)
decoder = Decoder(embedding_size, hidden_size, hidden_size, output_size)

seq2seq = Seq2Seq(encoder, decoder, 'cpu')
x_array = np.random.rand(batch, seq) * 10
x = torch.LongTensor(x_array)
out = seq2seq.forward(x)

expected_out = unit_test_values('seq2seq')
print('Close to out: ', expected_out.allclose(out, atol=1e-4))
```

```
Close to out: True
```

### 3.4: Train your Seq2Seq model

Now its time to combine what we have and train a Seq2Seq translator. We provided you with some training code and you can simply run them to see how your translator works. If you implemented everything correctly, you should see some meaningful translation in the output. You can modify the hyperparameters to improve the results. You can also tune the BATCH\_SIZE in section 1.2.

```
In [16]: # Hyperparameters. You are welcome to modify these
encoder_emb_size = 32
encoder_hidden_size = 64
encoder_dropout = 0.2

decoder_emb_size = 32
decoder_hidden_size = 64
decoder_dropout = 0.2

learning_rate = 1e-3
model_type = "LSTM"

EPOCHS = 30

#input size and output size
input_size = len(SRC.vocab)
output_size = len(TRG.vocab)
```

```
In [17]: # Declare models, optimizer, and loss function
encoder = Encoder(input_size, encoder_emb_size, encoder_hidden_size, decoder_hidden_size)
decoder = Decoder(decoder_emb_size, encoder_hidden_size, encoder_hidden_size, output_size)

seq2seq_model = Seq2Seq(encoder, decoder, device)
optimizer = optim.Adam(seq2seq_model.parameters(), lr = learning_rate, weight_decay=weight_decay)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer)
criterion = nn.CrossEntropyLoss(ignore_index=PAD_IDX)
```

```
In [18]: for epoch_idx in range(EPOCHS):
    print("-----")
    print("Epoch %d" % (epoch_idx+1))
    print("-----")

    train_loss, avg_train_loss = train(seq2seq_model, train_loader, optimizer, criterion, scheduler)
    scheduler.step(train_loss)

    val_loss, avg_val_loss = evaluate(seq2seq_model, valid_loader, criterion)

    avg_train_loss = avg_train_loss.item()
    avg_val_loss = avg_val_loss.item()
    print("Training Loss: %.4f. Validation Loss: %.4f. " % (avg_train_loss, avg_val_loss))
    print("Training Perplexity: %.4f. Validation Perplexity: %.4f. " % (np.exp(avg_train_loss), np.exp(avg_val_loss)))
```

```
-----
Epoch 1
-----
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/batch.py:23: UserWarning: Batch class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.
```

```
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.'.format(self.__class__.__name__), UserWarning)
```

```
Training Loss: 5.9377. Validation Loss: 5.2847.
Training Perplexity: 379.0475. Validation Perplexity: 197.2932.
```

```
-----
Epoch 2
-----
```

```
Training Loss: 5.2889. Validation Loss: 5.1208.
Training Perplexity: 198.1343. Validation Perplexity: 167.4769.
```

```
-----
Epoch 3
-----
```

```
Training Loss: 5.1809. Validation Loss: 5.0738.
Training Perplexity: 177.8425. Validation Perplexity: 159.7829.
```

```
-----
Epoch 4
-----
```

```
Training Loss: 5.1395. Validation Loss: 5.0451.
Training Perplexity: 170.6374. Validation Perplexity: 155.2558.
```

```
-----
Epoch 5
-----
```

```
Training Loss: 5.1113. Validation Loss: 5.0123.
```

Training Perplexity: 165.8916. Validation Perplexity: 150.2460.

Epoch 6

Training Loss: 5.0906. Validation Loss: 4.9949.

Training Perplexity: 162.4849. Validation Perplexity: 147.6548.

Epoch 7

Training Loss: 5.0564. Validation Loss: 4.9363.

Training Perplexity: 157.0180. Validation Perplexity: 139.2571.

Epoch 8

Training Loss: 4.9997. Validation Loss: 4.8827.

Training Perplexity: 148.3702. Validation Perplexity: 131.9863.

Epoch 9

Training Loss: 4.9525. Validation Loss: 4.8440.

Training Perplexity: 141.5282. Validation Perplexity: 126.9798.

Epoch 10

Training Loss: 4.9162. Validation Loss: 4.8114.

Training Perplexity: 136.4889. Validation Perplexity: 122.8975.

Epoch 11

Training Loss: 4.8832. Validation Loss: 4.7873.

Training Perplexity: 132.0471. Validation Perplexity: 119.9788.

Epoch 12

Training Loss: 4.8581. Validation Loss: 4.7676.

Training Perplexity: 128.7766. Validation Perplexity: 117.6384.

Epoch 13

Training Loss: 4.8370. Validation Loss: 4.7615.

Training Perplexity: 126.0933. Validation Perplexity: 116.9223.

Epoch 14

Training Loss: 4.8188. Validation Loss: 4.7402.

Training Perplexity: 123.8206. Validation Perplexity: 114.4550.

Epoch 15

-----  
Training Loss: 4.8013. Validation Loss: 4.7314.  
Training Perplexity: 121.6737. Validation Perplexity: 113.4599.  
-----

Epoch 16

-----  
Training Loss: 4.7844. Validation Loss: 4.7045.  
Training Perplexity: 119.6281. Validation Perplexity: 110.4411.  
-----

Epoch 17

-----  
Training Loss: 4.7641. Validation Loss: 4.6816.  
Training Perplexity: 117.2250. Validation Perplexity: 107.9404.  
-----

Epoch 18

-----  
Training Loss: 4.7438. Validation Loss: 4.6664.  
Training Perplexity: 114.8740. Validation Perplexity: 106.3144.  
-----

Epoch 19

-----  
Training Loss: 4.7227. Validation Loss: 4.6510.  
Training Perplexity: 112.4747. Validation Perplexity: 104.6911.  
-----

Epoch 20

-----  
Training Loss: 4.7014. Validation Loss: 4.6294.  
Training Perplexity: 110.0959. Validation Perplexity: 102.4544.  
-----

Epoch 21

-----  
Training Loss: 4.6828. Validation Loss: 4.6092.  
Training Perplexity: 108.0699. Validation Perplexity: 100.4022.  
-----

Epoch 22

-----  
Training Loss: 4.6614. Validation Loss: 4.5931.  
Training Perplexity: 105.7859. Validation Perplexity: 98.7976.  
-----

Epoch 23

-----  
Training Loss: 4.6422. Validation Loss: 4.5760.  
Training Perplexity: 103.7721. Validation Perplexity: 97.1252.  
-----

Epoch 24

Training Loss: 4.6216. Validation Loss: 4.5540.  
Training Perplexity: 101.6603. Validation Perplexity: 95.0124.

-----  
Epoch 25  
-----

Training Loss: 4.5999. Validation Loss: 4.5459.  
Training Perplexity: 99.4727. Validation Perplexity: 94.2442.

-----  
Epoch 26  
-----

Training Loss: 4.5801. Validation Loss: 4.5310.  
Training Perplexity: 97.5269. Validation Perplexity: 92.8491.

-----  
Epoch 27  
-----

Training Loss: 4.5600. Validation Loss: 4.5044.  
Training Perplexity: 95.5857. Validation Perplexity: 90.4115.

-----  
Epoch 28  
-----

Training Loss: 4.5394. Validation Loss: 4.4870.  
Training Perplexity: 93.6328. Validation Perplexity: 88.8531.

-----  
Epoch 29  
-----

Training Loss: 4.5162. Validation Loss: 4.4642.  
Training Perplexity: 91.4858. Validation Perplexity: 86.8530.

-----  
Epoch 30  
-----

Training Loss: 4.4945. Validation Loss: 4.4542.  
Training Perplexity: 89.5250. Validation Perplexity: 85.9847.

## 4: Train a Transformer

We will be implementing a one-layer Transformer **encoder** which, similar to an RNN, can encode a sequence of inputs and produce a final output of possibility of tokens in target language. This is the architecture:



You can refer to the [original paper](#) for more details. Before implementing a translator, we'll make a transformer binary classifier on the CoLA dataset below.

## The Corpus of Linguistic Acceptability (CoLA)

The Corpus of Linguistic Acceptability ([CoLA](#)) in its full form consists of 10657 sentences from 23 linguistics publications, expertly annotated for acceptability (grammaticality) by their original



authors. Native English speakers consistently report a sharp contrast in acceptability between pairs of sentences. Some examples include:

What did Betsy paint a picture of? (Correct)

What was a picture of painted by Betsy? (Incorrect)

You can read more info about the dataset [here](#). This is a binary classification task (predict 1 for correct grammar and 0 otherwise).

We will be using this dataset as a sanity checker for the forward pass of the Transformer architecture discussed in class. The general intuitive notion is that we will *encode* the sequence of tokens in the sentence, and then predict a binary output based on the final state that is the output of the model.

## Load the preprocessed data

We've appended a "CLS" token to the beginning of each sequence, which can be used to make predictions. The benefit of appending this token to the beginning of the sequence (rather than the end) is that we can extract it quite easily (we don't need to remove paddings and figure out the length of each individual sequence in the batch). We'll come back to this.

We've additionally already constructed a vocabulary and converted all of the strings of tokens into integers which can be used for vocabulary lookup for you. Feel free to explore the data here.

```
In [19]: train_inxs = np.load('./data/train_inxs.npy')
val_inxs = np.load('./data/val_inxs.npy')
train_labels = np.load('./data/train_labels.npy')
val_labels = np.load('./data/val_labels.npy')

# Load dictionary
word_to_ix = {}
with open("./data/word_to_ix.csv", "r") as f:
    reader = csv.reader(f)
    for line in reader:
        word_to_ix[line[0]] = line[1]
print("Vocabulary Size:", len(word_to_ix))

print(train_inxs.shape) # 7000 training instances, of (maximum/padded) Length 43 wor
print(val_inxs.shape) # 1551 validation instances, of (maximum/padded) Length 43 wor
print(train_labels.shape)
print(val_labels.shape)

d1 = torch.load('./data/d1.pt').to(device)
d2 = torch.load('./data/d2.pt').to(device)
d3 = torch.load('./data/d3.pt').to(device)
d4 = torch.load('./data/d4.pt').to(device)

Vocabulary Size: 1542
(7000, 43)
(1551, 43)
(7000,)
(1551,)
```

Instead of using numpy for this model, we will be using Pytorch to implement the forward pass. You will not need to implement the backward pass for the various layers in this assignment.

The file `models/Transformer.py` contains the model class and methods for each layer. This is where you will write your implementations.

## 4.1: Embeddings

We will format our input embeddings similarly to how they are constructed in [BERT \(source of figure\)](#). Recall from lecture that unlike a RNN, a Transformer does not include any positional information about the order in which the words in the sentence occur. Because of this, we need to append a positional encoding token at each position. (We will ignore the segment embeddings and [SEP] token here, since we are only encoding one sentence at a time). We have already appended the [CLS] token for you in the previous step.



Your first task is to implement the embedding lookup, including the addition of positional encodings. Open the file `Transformer.py` and complete all code parts for Deliverable 1.

```
In [22]: from models.Transformer import TransformerTranslator
inputs = train_inxs[0:2]
inputs = torch.LongTensor(inputs).to(device)

model = TransformerTranslator(input_size=len(word_to_ix), output_size=2, device=device)
model.to(device)

embeds = model.embed(inputs)

try:
    print("Difference:", torch.sum(torch.pairwise_distance(embeds, d1)).item()) # sh
except:
    print("NOT IMPLEMENTED")
```

Difference: 0.0017998493276536465

## 4.2: Multi-head Self-Attention

Attention can be computed in matrix-form using the following formula:



We want to have multiple self-attention operations, computed in parallel. Each of these is called a *head*. We concatenate the heads and multiply them with the matrix `attention_head_projection` to produce the output of this layer.

After every multi-head self-attention and feedforward layer, there is a residual connection + layer normalization. Make sure to implement this, using the following formula:



Open the file `models/transformer.py` and implement the `multihead_attention` function. We have already initialized all of the layers you will need in the constructor.

```
In [23]: hidden_states = model.multi_head_attention(embeds)

try:
    print("Difference:", torch.sum(torch.pairwise_distance(hidden_states, d2)).item())
except:
    print("NOT IMPLEMENTED")
```

Difference: 0.0017081360565498471

## 4.3: Element-Wise Feed-forward Layer

Open the file `models/transformer.py` and complete code for Deliverable 3 : the element-wise feed-forward layer consisting of two linear transformers with a ReLU layer in between.



```
In [24]: outputs = model.feedforward_layer(hidden_states)

try:
    print("Difference:", torch.sum(torch.pairwise_distance(outputs, d3)).item()) # s
except:
    print("NOT IMPLEMENTED")
```

Difference: 0.0017174689564853907

## 4.4: Final Layer

Open the file `models/transformer.py` and complete code for Deliverable 4 , to produce probability scores for all tokens in target language.

```
In [25]: scores = model.final_layer(outputs)

try:
    print("Difference:", torch.sum(torch.pairwise_distance(scores, d4)).item()) # sh
except:
    print("NOT IMPLEMENTED")
```

Difference: 2.628194124554284e-05

## 4.5: Putting it all together

Open the file `models/transformer.py` and complete the method `forward` , by putting together all of the methods you have developed in the right order to perform a full forward pass.

```
In [26]: inputs = train_inxs[0:2]
inputs = torch.LongTensor(inputs).to(device)

outputs = model.forward(inputs)

try:
    print("Difference:", torch.sum(torch.pairwise_distance(outputs, scores)).item())
except:
    print("NOT IMPLEMENTED")
```

Difference: 2.6229752620565705e-05

Great! We've just implemented a Transformer forward pass for translation. One of the big perks of using PyTorch is that with a simple training loop, we can rely on automatic differentiation ([autograd](#)) to do the work of the backward pass for us. This is not required for this assignment, but you can explore this on your own.

## 4.6: Train the Transformer

Now, we will extend your Transformer classifier to perform translation.

For simplicity, we have reduced the seq2seq translation problem to a token-for-token classification problem. The encoder will predict tokens of the target language. Hence, that this is *not* a real translation model like the seq2seq model from part 1. But as you will see, the Transformer can still learn to produce fluent text (with low perplexity). Also note that the BLEU score is a more typical metric for evaluating translations. Here, we use perplexity which is a language modeling metric for fluency.

Now you can start training the Transformer translator. We provided you with some training code and you can simply run them to see how your translator works. If you implemented everything correctly, you should see some meaningful translation in the output. Compare the results from the Seq2Seq model, which one is better? You can modify the hyperparameters to improve the results. You can also tune the BATCH\_SIZE in section 1.2.

```
In [27]: # Hyperparameters
learning_rate = 1e-2
EPOCHS = 10

# Model
trans_model = TransformerTranslator(input_size, output_size, device, max_length = MA

# optimizer = optim.Adam(model.parameters(), lr = learning_rate)
optimizer = torch.optim.Adam(trans_model.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer)
criterion = nn.CrossEntropyLoss(ignore_index=PAD_IDX)
```

```
In [28]: for epoch_idx in range(EPOCHS):
print("-----")
print("Epoch %d" % (epoch_idx+1))
print("-----")

train_loss, avg_train_loss = train(trans_model, train_loader, optimizer, criteri
scheduler.step(train_loss)

val_loss, avg_val_loss = evaluate(trans_model, valid_loader, criterion)

avg_train_loss = avg_train_loss.item()
avg_val_loss = avg_val_loss.item()
print("Training Loss: %.4f. Validation Loss: %.4f. " % (avg_train_loss, avg_val_
print("Training Perplexity: %.4f. Validation Perplexity: %.4f. " % (np.exp(avg_t
```

```
-----
Epoch 1
-----
```

```
/usr/local/lib/python3.7/dist-packages/torchtext/data/batch.py:23: UserWarning: Batc
h class will be retired soon and moved to torchtext.legacy. Please see the most rece
nt release notes for further information.
```

```
warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please
see the most recent release notes for further information.'.format(self.__class__.__
name__), UserWarning)
```

```
Training Loss: 4.2045. Validation Loss: 3.6560.
Training Perplexity: 66.9883. Validation Perplexity: 38.7070.
```

```
-----
Epoch 2
-----
```

```
Training Loss: 3.7553. Validation Loss: 3.5032.
Training Perplexity: 42.7456. Validation Perplexity: 33.2231.
-----
```

Epoch 3

-----  
Training Loss: 3.5786. Validation Loss: 3.4451.  
Training Perplexity: 35.8248. Validation Perplexity: 31.3459.  
-----

Epoch 4

-----  
Training Loss: 3.4495. Validation Loss: 3.3912.  
Training Perplexity: 31.4844. Validation Perplexity: 29.7027.  
-----

Epoch 5

-----  
Training Loss: 3.3392. Validation Loss: 3.3727.  
Training Perplexity: 28.1969. Validation Perplexity: 29.1560.  
-----

Epoch 6

-----  
Training Loss: 3.2676. Validation Loss: 3.3382.  
Training Perplexity: 26.2486. Validation Perplexity: 28.1695.  
-----

Epoch 7

-----  
Training Loss: 3.1795. Validation Loss: 3.3149.  
Training Perplexity: 24.0336. Validation Perplexity: 27.5210.  
-----

Epoch 8

-----  
Training Loss: 3.0992. Validation Loss: 3.3096.  
Training Perplexity: 22.1796. Validation Perplexity: 27.3747.  
-----

Epoch 9

-----  
Training Loss: 3.0280. Validation Loss: 3.2814.  
Training Perplexity: 20.6562. Validation Perplexity: 26.6118.  
-----

Epoch 10

-----  
Training Loss: 2.9833. Validation Loss: 3.3124.  
Training Perplexity: 19.7528. Validation Perplexity: 27.4496.  
-----

## Translations

Run the code below to see some of your translations. Modify to your liking.

```
In [29]: def translate(model, dataloader):  
        model.eval()  
        with torch.no_grad():  
            # Get the progress bar  
            progress_bar = tqdm(dataloader, ascii = True)
```

```

for batch_idx, data in enumerate(progress_bar):
    source = data.src.transpose(1,0)
    target = data.trg.transpose(1,0)

    translation = model(source)
    return target, translation

```

```

In [30]: # Select Transformer or Seq2Seq model
         # model = trans_model
         model = seq2seq_model

```

```

In [31]: #Set model equal to trans_model or seq2seq_model
         target, translation = translate(model, valid_loader)

```

```

0%|          | 0/8 [00:00<?, ?it/s]/usr/local/lib/python3.7/dist-packages/torchtext/data/batch.py:23: UserWarning: Batch class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.
  warnings.warn('{} class will be retired soon and moved to torchtext.legacy. Please see the most recent release notes for further information.'.format(self.__class__.__name__), UserWarning)
0%|          | 0/8 [00:00<?, ?it/s]

```

```

In [32]: raw = np.array([list(map(lambda x: TRG.vocab.itos[x], target[i])) for i in range(target.size())])
         print(raw)

```

```

[['<sos>' 'a' 'man' ... '<pad>' '<pad>' '<pad>']
 ['<sos>' 'a' 'man' ... '<pad>' '<pad>' '<pad>']
 ['<sos>' 'a' 'man' ... '<pad>' '<pad>' '<pad>']
 ...
 ['<sos>' 'boy' 'doing' ... '<pad>' '<pad>' '<pad>']
 ['<sos>' 'kids' 'are' ... '<pad>' '<pad>' '<pad>']
 ['<sos>' 'a' 'man' ... '<pad>' '<pad>' '<pad>']]

```

```

In [33]: token_trans = np.argmax(translation.cpu().numpy(), axis = 2)
         translated = np.array([list(map(lambda x: TRG.vocab.itos[x], token_trans[i])) for i in range(token_trans.size())])
         print(translated)

```

```

[['<unk>' 'a' 'man' ... '<eos>' '<eos>' '<eos>']
 ['<unk>' 'a' 'man' ... '<eos>' '<eos>' '<eos>']
 ['<unk>' 'a' 'man' ... '<eos>' '<eos>' '<eos>']
 ...
 ['<unk>' 'a' 'in' ... '<eos>' '<eos>' '<eos>']
 ['<unk>' 'people' 'are' ... '<eos>' '<eos>' '<eos>']
 ['<unk>' 'a' 'man' ... '<eos>' '<eos>' '<eos>']]

```